

Event detection and forecasting using KPGNNs on temporally dynamic graphs

improved with KL-divergence loss and Graph-ODE encoder

Ding, Yanna Gao, Suyi (Stanley) Saha, Bishwajit

Rensselaer Polytechnic Institute

- 1 Introduction to KPGNN
- 2 KL-divergence Loss
- 3 Graph ODE Encoder
- 4 Experiment and Empirical Result

Knowledge-Preserving Heterogeneous Graph Neural Network (KPGNN)

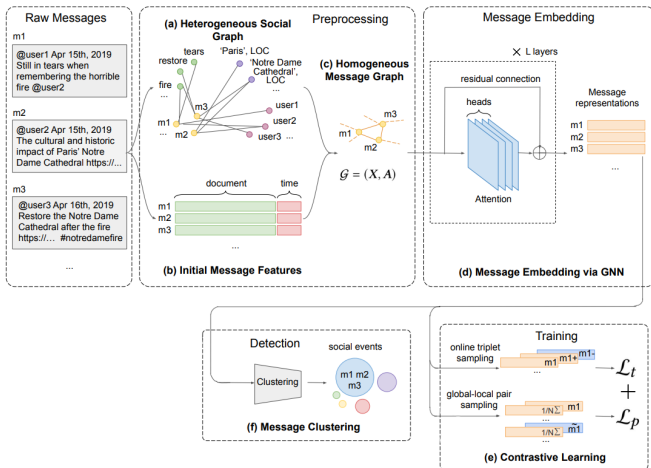


Figure 1: Workflow of KPGNN

KPGNN

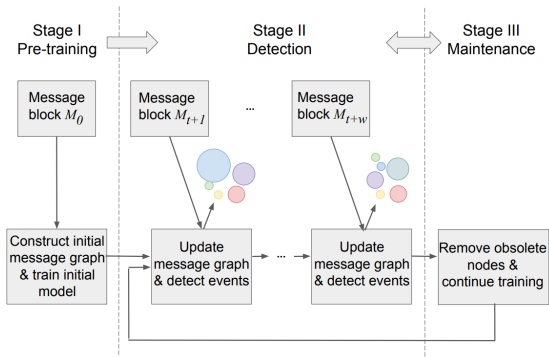


Figure 2: Training Stages of KPGNN

KPGNN vanilla

Algorithm 1: Model training in mini-batches (Stage 5)**Input:** Batch size β , number of batches B , number of layers L

```

1 function train( $\mathcal{G}, B, \beta, L$ )
2   for  $b = 1, 2, \dots, B$  do
3      $M^{(b)} = \{m_i\}_{i=1}^{\beta} \leftarrow$  sample messages from  $\mathcal{G}$ 
4     for  $l = 1, 2, \dots, L$  do
5        $\mathbf{h}_{m_i}^{(l)} \leftarrow \parallel_{\text{heads}} \left( \mathbf{h}_{m_i}^{(l-1)} \oplus \text{Aggregator}_{m_j \in \mathcal{N}(m_i)} \left( \text{Extractor} \left( \mathbf{h}_{m_j}^{(l-1)} \right) \right) \right)$ 
6        $\mathbf{h}_{m_i} \leftarrow \mathbf{h}_{m_i}^L$  among  $M^{(b)}$ 
7        $T = \left\{ (m_i, m_i^+, m_i^-) \right\}_{i=1}^{\beta} \leftarrow$  triplet sampling from  $M^{(b)}$ 
8        $\mathcal{L}_t \leftarrow$  triplet loss
9        $\mathbf{s}, \tilde{\mathbf{h}}_{m_i} \leftarrow$  calculate summary and corrupted representation
10       $\mathcal{L}_p \leftarrow$  cross entropy loss
11      Back-propagation to update parameters

```

KPGNN improved with GODE and KL divergence loss

Algorithm 2: Model training in mini-batches (Stage 5)**Input:** Batch size β , number of batches B , number of layers L

```

1 function train( $\mathcal{G}, B, \beta, L$ )
2   for  $b = 1, 2, \dots, B$  do
3      $M^{(b)} = \{m_i\}_{i=1}^{\beta} \leftarrow$  sample messages from  $\mathcal{G}$ 
4     for  $l = 1, 2, \dots, L$  do
5        $\mathbf{h}_{m_i}^{(l)} \leftarrow$  (new) GODE encoder
6      $\mathbf{h}_{m_i} \leftarrow \mathbf{h}_{m_i}^L$  among  $M^{(b)}$ 
7      $T = \{(m_i, m_i^+, m_i^-)\}_{i=1}^{\beta} \leftarrow$  triplet sampling from  $M^{(b)}$ 
8      $\mathcal{L}_t \leftarrow$  triplet loss
9      $\mathbf{s}, \tilde{\mathbf{h}}_{m_i} \leftarrow$  calculate summary and corrupted representation
10     $\mathcal{L}_p \leftarrow$  cross entropy loss
11     $\mathcal{L}_{kl} \leftarrow$  KL divergence loss
12    Back-propagation to update parameters

```

- 1 Introduction to KPGNN
- 2 KL-divergence Loss
- 3 Graph ODE Encoder
- 4 Experiment and Empirical Result

KL-divergence Loss

KL-divergence loss, also known as Kullback-Leibler divergence loss is used to measure the distance between two distributions. By including this loss in the model, we can obtain better clustering results.

$$\mathcal{L}_{kl} = \text{KL}(P\|Q) = \sum_{ij} p_{ij} \log \left(\frac{p_{ij}}{q_{ij}} \right)$$

Predicted soft assignment:

$$q_{ij} = \frac{\left(1 + \frac{1}{\alpha} \|z_i - \mu_j\|^2\right)^{-\frac{1}{2}(\alpha+1)}}{\sum_{j'} \left(1 + \frac{1}{\alpha} \|z_i - \mu_{j'}\|^2\right)^{-\frac{1}{2}(\alpha+1)}}$$

Target distribution:

$$p_{ij} = \frac{q_{ij}^2 / \sum_i q_{ij}}{\sum_{j'} (q_{ij'}^2 / \sum_i q_{ij'})}$$

KL-divergence Loss

- For p_i , a naive approach would be setting it to a delta distribution (to the nearest centroid) for data points above a confidence threshold and ignoring the rest.
- However, because q_i is soft assignment, it is more flexible to use softer probabilistic targets.
- Here p_i is defined by q_i and the intent of optimization is to match Q with P , thus we can regard it as a kind of self-training.
- Strengthen the predictions by putting more importance on data points assigned with high confidence that eventually improves cluster purity.
- Normalize loss contribution of each centroid to prevent large clusters from distorting the hidden feature space.

- 1 Introduction to KPGNN
- 2 KL-divergence Loss
- 3 Graph ODE Encoder
- 4 Experiment and Empirical Result

Graph ODE Encoder

We replace the original encoder with a continuous version of layered graph neural networks and refer it to the Graph Ordinary Differential Equation (GODE) model. The embedding \mathbf{h}_{m_i} for node i is obtained by solving an initial value problem (IVP) for a coupled ODE system.

$$\mathbf{h}_{m_i}^{(t_1)} = \mathbf{h}_{m_i}^{(t_0)} + \int_{t_0}^{t_1} \text{Aggregator}_{m_j \in \mathcal{N}(m_i)}(\text{Extractor}(\mathbf{h}_{m_i}(t))) dt$$

where $\mathbf{h}_{m_i}^{(t_0)}$ is the initial condition of the hidden vector.

The corresponding discrete transformation is

$$\mathbf{h}_{m_i}^{(t)} = \mathbf{h}_{m_i}^{(t-1)} + \text{Aggregator}_{m_j \in \mathcal{N}(m_i)}(\text{Extractor}(\mathbf{h}_{m_j}^{(t-1)}))$$

Graph ODE Encoder: Expressivity

We use an example from the Neural Ordinary Differential Equations model to show the expressivity of ODE-based neural networks.

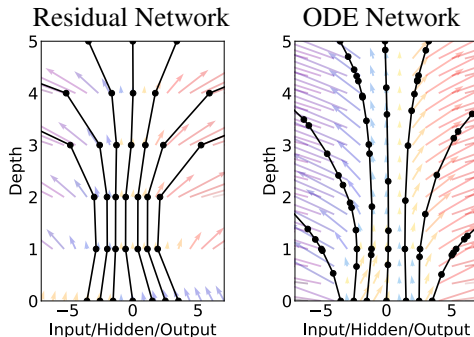


Figure 3: Neural ordinary differential equations, by Chen et al. The residual network has an update rule of $h_{t+1} = h_t + f(h_t)$, and the ODE network defines the final embedding as the solution to an IVP $h = h^{(0)} + \int_{t_0}^{t_1} f(h(t), t) dt$.

Graph ODE Encoder

We use the following instance of the above model. The expression for the hidden state of message m_i can be written as:

$$\mathbf{h}_{m_i}^{(T)} = \mathbf{h}_{m_i}^{(0)} + \frac{1}{\sqrt{k_i}} \int_0^T \sigma \left[\sum_{m_v \in \mathcal{N}(m_i)} \frac{1}{\sqrt{k_v}} (W \mathbf{h}_{m_v}(t) + b) \right] dt, \quad \mathbf{h}_{m_i}^{(0)} = f_{\text{enc}}(\mathbf{x}_{m_i})$$

The ODE is numerically solved using Runge-kutta4 (Dormand & Prince (1980)).

- 1 Introduction to KPGNN
- 2 KL-divergence Loss
- 3 Graph ODE Encoder
- 4 Experiment and Empirical Result**

Experiment Setup

We conduct experiments on a subset of the Twitter dataset. Message block 0 includes 500 messages. Blocks 1-21 contains 100 messages and the corresponding heterogeneous knowledge representation graph consist of 100 nodes.

We adopt the *latest message strategy* to train the model. In the initial stage, the model is trained on message block M_0 . For blocks $i = 1, \dots, 21$, we first perform inference to get the test NMI for block i and if $i = 0 \bmod \text{window size}$, we train T iterations on block i .

NMI is a metric commonly used to calculate the clustering quality between the predicted labels and ground truth labels. A value of 1 indicates perfect clustering while a value of 0 indicates completely wrong class labels. Higher is better.

Test Result: KL-divergence




Figure 4: KPGNN vs KPGNN with kld loss at training stage.

Test Result: KL-divergence

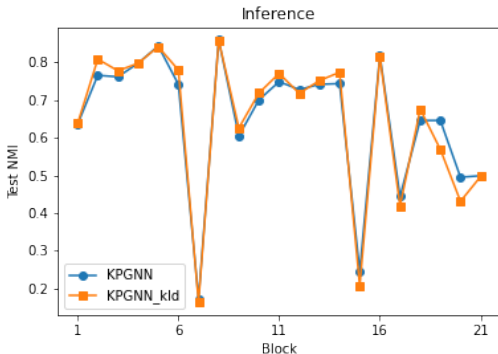


Figure 5: KPGNN vs KPGNN with kld loss at inference stage.

Test Result: GODE




Figure 6: Test NMI for KPGNN and GODE.

Test Result: GODE




Figure 7: GODE with different hyperparameters. We fix the batch size at 150 and vary window size from $\{1, 3, 5, 7, 9\}$ and the embedding dimension from $\{16, 32, 48, 64, 80\}$.

Test Result: KL-divergence + GODE




Figure 8: Test NMI for KPGNN, GODE, and the combination of GODE and KL-divergence.

References

Thank you for listening!

- Yuwei Cao, Hao Peng, and Philip S Yu. Multi-information source hin for medical concept embedding. In Advances in Knowledge Discovery and Data Mining: 24th Pacific-Asia Conference, PAKDD 2020, Singapore, May 11–14, 2020, Proceedings, Part II 24, pp. 396–408. Springer, 2020.
- Yuwei Cao, Hao Peng, Jia Wu, Yingdong Dou, Jianxin Li, and Philip S Yu. Knowledge-preserving incremental social event detection via heterogeneous gnn. In Proceedings of the Web Conference 2021, pp. 3383–3395, 2021.
- John R Dormand and Peter J Prince. A family of embedded runge-kutta formulae. Journal of computational and applied mathematics, 6(1):19–26, 1980.
- Andrew J McMin, Yashar Moshfeghi, and Joemon M Jose. Building a large-scale corpus for evaluating event detection on twitter. In Proceedings of the 22nd ACM international conference on Information & Knowledge Management, pp. 409–418, 2013.
- Hao Peng, Jianxin Li, Qiran Gong, Yangqiu Song, Yuanxing Ning, Kunfeng Lai, and Philip S Yu. Fine-grained event categorization with heterogeneous graph convolutional networks. arXiv preprint arXiv:1906.04580, 2019.
- Michael Poli, Stefano Massaroli, Junyoung Park, Atsushi Yamashita, Hajime Asama, and Jinkyoo Park. Graph neural ordinary differential equations. arXiv preprint arXiv:1911.07532, 2019.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention is all you need, 2017.
- Junyuan Xie, Ross Girshick, and Ali Farhadi. Unsupervised deep embedding for clustering analysis. In International conference on machine learning, pp. 478–487. PMLR, 2016.